

Package ‘xgboost’

December 22, 2014

Type Package

Title eXtreme Gradient Boosting

Version 0.3-2

Date 2014-08-23

Author Tianqi Chen <tianqi.tchen@gmail.com>, Tong He <hetong007@gmail.com>

Maintainer Tong He <hetong007@gmail.com>

Description This package is a R wrapper of xgboost, which is short for eXtreme Gradient Boosting. It is an efficient and scalable implementation of gradient boosting framework. The package includes efficient linear model solver and tree learning algorithms. The package can automatically do parallel computation with OpenMP, and it can be more than 10 times faster than existing gradient boosting packages such as gbm. It supports various objective functions, including regression, classification and ranking. The package is made to be extensible, so that users are also allowed to define their own objectives easily.

License Apache License (== 2.0) | file LICENSE

URL <https://github.com/tqchen/xgboost>

BugReports <https://github.com/tqchen/xgboost/issues>

Depends R (>= 2.10)

Imports Matrix (>= 1.1-0), methods

NeedsCompilation yes

Repository CRAN

Date/Publication 2014-09-07 21:54:44

R topics documented:

agaricus.test	2
agaricus.train	3
getinfo	3
predict,xgb.Booster-method	4

setinfo	5
slice	6
xgb.cv	6
xgb.DMatrix	8
xgb.DMatrix.save	8
xgb.dump	9
xgb.load	10
xgb.save	10
xgb.train	11
xgboost	12
Index	14

agaricus.test	<i>Test part from Mushroom Data Set</i>
---------------	---

Description

This data set is originally from the Mushroom data set, UCI Machine Learning Repository.

Usage

```
data(agaricus.test)
```

Format

A list containing a label vector, and a dgCMatrix object with 1611 rows and 127 variables

Details

This data set includes the following fields:

- label the label for each record
- data a sparse Matrix of dgCMatrix class, with 127 columns.

References

<https://archive.ics.uci.edu/ml/datasets/Mushroom>
Bache, K. & Lichman, M. (2013). UCI Machine Learning Repository [<http://archive.ics.uci.edu/ml>].
Irvine, CA: University of California, School of Information and Computer Science.

agaricus.train	<i>Training part from Mushroom Data Set</i>
----------------	---

Description

This data set is originally from the Mushroom data set, UCI Machine Learning Repository.

Usage

```
data(agaricus.train)
```

Format

A list containing a label vector, and a dgCMatrix object with 6513 rows and 127 variables

Details

This data set includes the following fields:

- label the label for each record
- data a sparse Matrix of dgCMatrix class, with 127 columns.

References

<https://archive.ics.uci.edu/ml/datasets/Mushroom>

Bache, K. & Lichman, M. (2013). UCI Machine Learning Repository [<http://archive.ics.uci.edu/ml>]. Irvine, CA: University of California, School of Information and Computer Science.

getinfo	<i>Get information of an xgb.DMatrix object</i>
---------	---

Description

Get information of an xgb.DMatrix object

Usage

```
getinfo(object, ...)
```

```
## S4 method for signature 'xgb.DMatrix'  
getinfo(object, name)
```

Arguments

object	Object of class "xgb.DMatrix"
name	the name of the field to get
...	other parameters

Examples

```
data(agaricus.train, package='xgboost')
train <- agaricus.train
dtrain <- xgb.DMatrix(train$data, label=train$label)
labels <- getinfo(dtrain, 'label')
setinfo(dtrain, 'label', 1-labels)
labels2 <- getinfo(dtrain, 'label')
stopifnot(all(labels2 == 1-labels))
```

predict,xgb.Booster-method

Predict method for eXtreme Gradient Boosting model

Description

Predicted values based on xgboost model object.

Usage

```
## S4 method for signature 'xgb.Booster'
predict(object, newdata, outputmargin = FALSE,
        ntreetlimit = NULL)
```

Arguments

object	Object of class "xgb.Boost"
newdata	takes matrix, dgCMatix, local data file or xgb.DMatrix.
outputmargin	whether the prediction should be shown in the original value of sum of functions, when outputmargin=TRUE, the prediction is untransformed margin value. In logistic regression, outputmargin=T will output value before logistic transformation.
ntreetlimit	limit number of trees used in prediction, this parameter is only valid for gbtrees, but not for gblinears. set it to be value bigger than 0. It will use all trees by default.

Examples

```
data(agaricus.train, package='xgboost')
data(agaricus.test, package='xgboost')
train <- agaricus.train
test <- agaricus.test
bst <- xgboost(data = train$data, label = train$label, max.depth = 2,
              eta = 1, nround = 2, objective = "binary:logistic")
pred <- predict(bst, test$data)
```

setinfo

Set information of an xgb.DMatrix object

Description

Set information of an xgb.DMatrix object

Usage

```
setinfo(object, ...)

## S4 method for signature 'xgb.DMatrix'
setinfo(object, name, info)
```

Arguments

object	Object of class "xgb.DMatrix"
name	the name of the field to get
info	the specific field of information to set
...	other parameters

Examples

```
data(agaricus.train, package='xgboost')
train <- agaricus.train
dtrain <- xgb.DMatrix(train$data, label=train$label)
labels <- getinfo(dtrain, 'label')
setinfo(dtrain, 'label', 1-labels)
labels2 <- getinfo(dtrain, 'label')
stopifnot(all(labels2 == 1-labels))
```

slice	<i>Get a new DMatrix containing the specified rows of original xgb.DMatrix object</i>
-------	---

Description

Get a new DMatrix containing the specified rows of original xgb.DMatrix object

Usage

```
slice(object, ...)

## S4 method for signature 'xgb.DMatrix'
slice(object, idxset, ...)
```

Arguments

object	Object of class "xgb.DMatrix"
idxset	a integer vector of indices of rows needed
...	other parameters

Examples

```
data(agaricus.train, package='xgboost')
train <- agaricus.train
dtrain <- xgb.DMatrix(train$data, label=train$label)
dsub <- slice(dtrain, 1:3)
```

xgb.cv	<i>Cross Validation</i>
--------	-------------------------

Description

The cross valudation function of xgboost

Usage

```
xgb.cv(params = list(), data, nrounds, nfold, label = NULL, showsd = TRUE,
        metrics = list(), obj = NULL, feval = NULL, ...)
```

Arguments

params	<p>the list of parameters. Commonly used ones are:</p> <ul style="list-style-type: none"> • objective objective function, common ones are <ul style="list-style-type: none"> – reg:linear linear regression – binary:logistic logistic regression for classification • eta step size of each boosting step • max.depth maximum depth of the tree • nthread number of thread used in training, if not set, all threads are used <p>See https://github.com/tqchen/xgboost/wiki/Parameters for further details. See also demo/ for walkthrough example in R.</p>
data	takes an xgb.DMatrix as the input.
nrounds	the max number of iterations
nfold	number of folds used
label	option field, when data is Matrix
showsd	boolean, whether show standard deviation of cross validation
metrics,	<p>list of evaluation metrics to be used in corss validation, when it is not specified, the evaluation metric is chosen according to objective function. Possible options are:</p> <ul style="list-style-type: none"> • error binary classification error rate • rmse Rooted mean square error • logloss negative log-likelihood function • auc Area under curve • merror Exact matching error, used to evaluate multi-class classification
obj	customized objective function. Returns gradient and second order gradient with given prediction and dtrain,
feval	customized evaluation function. Returns list(metric='metric-name', value='metric-value') with given prediction and dtrain,
...	other parameters to pass to params.

Details

This is the cross validation function for xgboost

Parallelization is automatically enabled if OpenMP is present. Number of threads can also be manually specified via "nthread" parameter.

This function only accepts an xgb.DMatrix object as the input.

Examples

```
data(agaricus.train, package='xgboost')
dtrain <- xgb.DMatrix(agaricus.train$data, label = agaricus.train$label)
history <- xgb.cv(data = dtrain, nround=3, nfold = 5, metrics=list("rmse","auc"),
                  "max.depth"=3, "eta"=1, "objective"="binary:logistic")
```

xgb.DMatrix	<i>Construct xgb.DMatrix object</i>
-------------	-------------------------------------

Description

Construct xgb.DMatrix object from dense matrix, sparse matrix or local file.

Usage

```
xgb.DMatrix(data, info = list(), missing = 0, ...)
```

Arguments

data	a matrix object, a dgCMatrx object or a character indicating the data file.
info	a list of information of the xgb.DMatrix object
missing	Missing is only used when input is dense matrix, pick a float
...	other information to pass to info.

Examples

```
data(agaricus.train, package='xgboost')
train <- agaricus.train
dtrain <- xgb.DMatrix(train$data, label=train$label)
xgb.DMatrix.save(dtrain, 'xgb.DMatrix.data')
dtrain <- xgb.DMatrix('xgb.DMatrix.data')
```

xgb.DMatrix.save	<i>Save xgb.DMatrix object to binary file</i>
------------------	---

Description

Save xgb.DMatrix object to binary file

Usage

```
xgb.DMatrix.save(DMatrix, fname)
```

Arguments

DMatrix	the DMatrix object
fname	the name of the binary file.

Examples

```
data(agaricus.train, package='xgboost')
train <- agaricus.train
dtrain <- xgb.DMatrix(train$data, label=train$label)
xgb.DMatrix.save(dtrain, 'xgb.DMatrix.data')
dtrain <- xgb.DMatrix('xgb.DMatrix.data')
```

xgb.dump	<i>Save xgboost model to text file</i>
----------	--

Description

Save a xgboost model to text file. Could be parsed later.

Usage

```
xgb.dump(model, fname, fmap = "")
```

Arguments

model	the model object.
fname	the name of the binary file.
fmap	feature map file representing the type of feature. Detailed description could be found at https://github.com/tqchen/xgboost/wiki/Binary-Classification#dump-model . See demo/ for walkthrough example in R, and https://github.com/tqchen/xgboost/blob/master/demo/data/featmap.txt for example Format.

Examples

```
data(agaricus.train, package='xgboost')
data(agaricus.test, package='xgboost')
train <- agaricus.train
test <- agaricus.test
bst <- xgboost(data = train$data, label = train$label, max.depth = 2,
              eta = 1, nround = 2, objective = "binary:logistic")
xgb.dump(bst, 'xgb.model.dump')
```

xgb.load	<i>Load xgboost model from binary file</i>
----------	--

Description

Load xgboost model from the binary model file

Usage

```
xgb.load(modelfile)
```

Arguments

modelfile	the name of the binary file.
-----------	------------------------------

Examples

```
data(agaricus.train, package='xgboost')
data(agaricus.test, package='xgboost')
train <- agaricus.train
test <- agaricus.test
bst <- xgboost(data = train$data, label = train$label, max.depth = 2,
              eta = 1, nround = 2, objective = "binary:logistic")
xgb.save(bst, 'xgb.model')
bst <- xgb.load('xgb.model')
pred <- predict(bst, test$data)
```

xgb.save	<i>Save xgboost model to binary file</i>
----------	--

Description

Save xgboost model from xgboost or xgb.train

Usage

```
xgb.save(model, fname)
```

Arguments

model	the model object.
fname	the name of the binary file.

Examples

```
data(agaricus.train, package='xgboost')
data(agaricus.test, package='xgboost')
train <- agaricus.train
test <- agaricus.test
bst <- xgboost(data = train$data, label = train$label, max.depth = 2,
              eta = 1, nround = 2, objective = "binary:logistic")
xgb.save(bst, 'xgb.model')
bst <- xgb.load('xgb.model')
pred <- predict(bst, test$data)
```

xgb.train

eXtreme Gradient Boosting Training

Description

The training function of xgboost

Usage

```
xgb.train(params = list(), data, nrounds, watchlist = list(), obj = NULL,
          feval = NULL, verbose = 1, ...)
```

Arguments

params	<p>the list of parameters. Commonly used ones are:</p> <ul style="list-style-type: none"> • objective objective function, common ones are <ul style="list-style-type: none"> – reg:linear linear regression – binary:logistic logistic regression for classification • eta step size of each boosting step • max.depth maximum depth of the tree • nthread number of thread used in training, if not set, all threads are used <p>See https://github.com/tqchen/xgboost/wiki/Parameters for further details. See also demo/ for walkthrough example in R.</p>
data	takes an xgb.DMatrix as the input.
nrounds	the max number of iterations
watchlist	what information should be printed when verbose=1 or verbose=2. Watchlist is used to specify validation set monitoring during training. For example user can specify watchlist=list(validation1=mat1, validation2=mat2) to watch the performance of each round's model on mat1 and mat2
obj	customized objective function. Returns gradient and second order gradient with given prediction and dtrain,
feval	customized evaluation function. Returns list(metric='metric-name', value='metric-value') with given prediction and dtrain,
verbose	If 0, xgboost will stay silent. If 1, xgboost will print information of performance. If 2, xgboost will print information of both
...	other parameters to pass to params.

Details

This is the training function for xgboost.

Parallelization is automatically enabled if OpenMP is present. Number of threads can also be manually specified via "nthread" parameter.

This function only accepts an xgb.DMatrix object as the input. It supports advanced features such as watchlist, customized objective function, therefore it is more flexible than [xgboost](#).

Examples

```
data(agaricus.train, package='xgboost')
dtrain <- xgb.DMatrix(agaricus.train$data, label = agaricus.train$label)
dtest <- dtrain
watchlist <- list(eval = dtest, train = dtrain)
param <- list(max.depth = 2, eta = 1, silent = 1)
logregobj <- function(preds, dtrain) {
  labels <- getinfo(dtrain, "label")
  preds <- 1/(1 + exp(-preds))
  grad <- preds - labels
  hess <- preds * (1 - preds)
  return(list(grad = grad, hess = hess))
}
evalerror <- function(preds, dtrain) {
  labels <- getinfo(dtrain, "label")
  err <- as.numeric(sum(labels != (preds > 0)))/length(labels)
  return(list(metric = "error", value = err))
}
bst <- xgb.train(param, dtrain, nround = 2, watchlist, logregobj, evalerror)
```

xgboost

eXtreme Gradient Boosting (Tree) library

Description

A simple interface for xgboost in R

Usage

```
xgboost(data = NULL, label = NULL, params = list(), nrounds,
        verbose = 1, ...)
```

Arguments

data	takes matrix, dgCMatix, local data file or xgb.DMatrix.
label	the response variable. User should not set this field,
params	the list of parameters. Commonly used ones are: <ul style="list-style-type: none"> • objective objective function, common ones are <ul style="list-style-type: none"> – reg:linear linear regression

- `binary:logistic` logistic regression for classification

- `eta` step size of each boosting step
- `max.depth` maximum depth of the tree
- `nthread` number of thread used in training, if not set, all threads are used

See <https://github.com/tqchen/xgboost/wiki/Parameters> for further details. See also `demo/` for walkthrough example in R.

<code>nrounds</code>	the max number of iterations
<code>verbose</code>	If 0, xgboost will stay silent. If 1, xgboost will print information of performance. If 2, xgboost will print information of both performance and construction progress information
<code>...</code>	other parameters to pass to <code>params</code> .

Details

This is the modeling function for xgboost.

Parallelization is automatically enabled if OpenMP is present. Number of threads can also be manually specified via `"nthread"` parameter

Examples

```
data(agaricus.train, package='xgboost')
data(agaricus.test, package='xgboost')
train <- agaricus.train
test <- agaricus.test
bst <- xgboost(data = train$data, label = train$label, max.depth = 2,
              eta = 1, nround = 2, objective = "binary:logistic")
pred <- predict(bst, test$data)
```

Index

*Topic **datasets**

agaricus.test, [2](#)

agaricus.train, [3](#)

agaricus.test, [2](#)

agaricus.train, [3](#)

getinfo, [3](#)

getinfo, xgb.DMatrix-method (getinfo), [3](#)

predict, xgb.Booster-method, [4](#)

setinfo, [5](#)

setinfo, xgb.DMatrix-method (setinfo), [5](#)

slice, [6](#)

slice, xgb.DMatrix-method (slice), [6](#)

xgb.cv, [6](#)

xgb.DMatrix, [8](#)

xgb.DMatrix.save, [8](#)

xgb.dump, [9](#)

xgb.load, [10](#)

xgb.save, [10](#)

xgb.train, [11](#)

xgboost, [12](#), [12](#)